Quantum-Enhanced Active Learning for Accelerated Materials Discovery

Arnav Kapoor

Indian Institute of Science Education and Research, Bhopal

Bhopal, India

arnavkapoor23@iiserb.ac.in

Abstract—This paper develops a quantum-enhanced active learning framework for materials discovery by combining quantum-inspired state encoding with multi-observable uncertainty quantification. The approach derives uncertainty metrics by capturing coupled structure-property relationships. Evaluated on band gap and formation energy prediction tasks in five independent runs comparing nine baselines, the quantum-enhanced method consistently achieves a higher test R² and improved sample efficiency, reducing the experimental budget by 25-35% (paired t-tests, $p < 0.01$).

Keywords: quantum-inspired algorithms, active learning, materials discovery, uncertainty quantification, machine learning

I. Introduction

Accelerating material discovery is critical to address energy storage and catalysis challenges. Active learning prioritizes informative experiments and reducing evaluations needed to identify promising candidates [1], [2].

Classical active learning relies on scalar uncertainty estimates that miss coupled phenomena in material systems. For instance, band gap and formation energy correlate through structural factors; selecting samples uncertain in one but certain in the other misses opportunities to reduce both simultaneously. This coupling problem is severe in multi-property design spaces [1], [3].

Quantum mechanics offers a conceptual framework for coupled systems through superposition and entanglement. Quantum-inspired classical algorithms extract these principles—amplitude-based feature representations, multiple non-commuting observables, coupled uncertainty metrics—without requiring quantum hardware [4].

This paper develops a quantum-enhanced active learning framework for materials discovery with: (i) quantum-inspired state encoding of materials features; (ii) multi-observable framework with cross-observable covariance computation; (iii) uncertainty aggregation combining variances and covariances; (iv) model-agnostic active learning loop; (v) validation on band gap/formation energy (5 trials, 9 baselines, $p < 0.01$).

II. Related Work

Active Learning and Batch Selection: Classical active learning for materials design commonly uses uncertainty sampling, Query-by-Committee, expected improvement, diversity/coverage heuristics, and core-set style batch selection [1]–[3], [5], [6]. These strategies rely on scalar uncertainty, which can miss coupled structure-property effects.

Uncertainty Estimation: Bayesian approximations such as dropout and deep ensembles improve predictive uncertainty but still aggregate properties into single scores [7], [8]. This approach instead couples multiple observables with covariance terms to retain cross-property structure.

Quantum-Inspired and Quantum Machine Learning: Quantum-inspired classical methods use amplitude encodings and linear-algebraic speedups to represent complex correlations [9], [10]. Variational and NISQ-era quantum machine learning advances motivate multi-observable formulations and structured uncertainty modeling [11]–[13]. The method leverages these ideas while remaining hardware-agnostic.

III. Proposed Method

The proposed quantum-enhanced active learning framework integrates quantum-inspired representations with modern machine learning to improve materials discovery efficiency [11]–[13]. This section details the mathematical formulation, uncertainty quantification strategy, and algorithmic implementation.

Quantum State Representation: Each candidate material is encoded as a quantum state $|\psi_i\rangle$ in a d-dimensional Hilbert space spanned by materials features. The state preparation maps normalized feature vectors to amplitude representations:

$$|\psi_i\rangle = \sum_{j=1}^{d} \alpha_{ij}|f_j\rangle, \tag{1}$$

where $\{|f_j\rangle\}_{j=1}^{d}$ form an orthonormal basis corresponding to normalized features (atomic radius, electronegativity, formation energy, coordination number, etc.), and $\alpha_{ij} \in \mathbb{C}$ are complex amplitudes satisfying normalization $\sum_j |\alpha_{ij}|^2 = 1$. The amplitudes are constructed via a phase-encoding scheme: $\alpha_{ij} = \sqrt{p_{ij}}e^{i\phi_{ij}}$, where p_{ij} are probabilities derived from normalized feature values and ϕ_{ij} encode feature correlations. This amplitude encoding (Eq. (1)) maps materials features into a normalized quantum-like representation [9].

Multi-Observable Framework: To capture different facets of materials properties, three quantum observables are defined

as Hermitian operators acting on the feature space [14], [15]:

$$\hat{O}_{\text{structure}} = \sum_{i,j} w_{ij}^{\text{struct}} |f_i\rangle\langle f_j|, \qquad (2)$$

$$\hat{O}_{\text{electronic}} = \sum_{i,j} w_{ij}^{\text{elec}} |f_i\rangle\langle f_j|, \qquad (3)$$

$$\hat{O}_{\text{thermodynamic}} = \sum_{i,j} w_{ij}^{\text{thermo}} |f_i\rangle\langle f_j|, \qquad (4)$$

where $w_{ij}^{(\cdot)}$ are weights initialized from domain knowledge. Non-commuting observables ($[\hat{O}_k, \hat{O}_l] \neq 0$) capture quantum-like uncertainty relations. *Construction guidelines:* (1) Identify 3-4 complementary property aspects; (2) Map to feature subsets (structural: geometry, electronic: band structure, thermodynamic: stability); (3) Initialize w_{ij} as block-diagonal or full, normalized to unit spectral norm; (4) Ensure diversity—redundant types degrade performance $\sim 3\%$; (5) Normalize variance scales; (6) Optionally tune via gradient descent.

Uncertainty Quantification: For each observable $\hat{O}_k$, the quantum variance quantifies intrinsic uncertainty (Eq. (5)):

$$\sigma^2(\hat{O}_k) = \langle\psi|\hat{O}_k^2|\psi\rangle - \langle\psi|\hat{O}_k|\psi\rangle^2. \qquad (5)$$

Cross-observable covariances capture correlations (Eq. (6)):

$$\text{Cov}(\hat{O}_k, \hat{O}_l) = \langle\psi|\frac{\hat{O}_k\hat{O}_l + \hat{O}_l\hat{O}_k}{2}|\psi\rangle - \langle\psi|\hat{O}_k|\psi\rangle\langle\psi|\hat{O}_l|\psi\rangle. \qquad (6)$$

The symmetrized product ensures real-valued covariance. These terms encode coupled uncertainties across property domains.

Selection Strategy: The total uncertainty aggregates individual variances and pairwise covariances using complex-valued combination coefficients [7], [8]:

$$U_{\text{total}} = \sqrt{\sum_k |\alpha_k|^2 \sigma^2(\hat{O}_k) + \sum_{k \neq l} \text{Re}(\alpha_k^* \alpha_l) \text{Cov}(\hat{O}_k, \hat{O}_l)}, \qquad (7)$$

where $\alpha_k \in \mathbb{C}$ are complex coefficients tuning the relative contribution of each observable, and $\text{Re}(\cdot)$ denotes the real part. The coefficients $\{\alpha_k\}$ can be learned via cross-validation or set based on domain priorities (e.g., higher weight for electronic properties in band gap prediction). Candidates are ranked by U_{total} (Eq. (7)) in descending order, and the top b candidates form the query batch for experimental evaluation.

Computational Complexity: The framework computes expectations as $\langle\psi|\hat{O}|\psi\rangle = \sum_{ij} \alpha_i^* w_{ij} \alpha_j$, which requires $\mathcal{O}(d^2)$ operations per observable if the weight matrix w is dense. For sparse or structured observables (e.g., banded or low-rank), this reduces to $\mathcal{O}(d)$ or $\mathcal{O}(rd)$ where r is the rank. Computing uncertainties for N candidates and K observables thus requires $\mathcal{O}(NKd^2)$ in the worst case, but structured operators and efficient linear algebra (matrix-vector products, Cholesky decompositions) yield practical complexity $\mathcal{O}(NKd)$, making the approach scalable to thousands of candidates with typical feature dimensions $d \sim 50\text{-}200$.

Active Learning Loop: The overall procedure is summarized in Alg. 1. Starting from an initial labeled set $\mathcal{D}_0$ of size

---

Algorithm 1 Quantum-Enhanced Active Learning

Input: Candidate materials $\{M_i\}_{i=1}^N$, initial labeled set $\mathcal{D}_0$, batch size b
Output: Trained model f and expanded labeled set $\mathcal{D}$
Initialize quantum states $\{|\psi_i\rangle\}$ for all candidates via feature encoding
Define quantum observables $\{\hat{O}_k\}_{k=1}^K$ with domain-informed weights
Set combination coefficients $\{\alpha_k\}_{k=1}^K$
$t \leftarrow 0, \mathcal{D}_t \leftarrow \mathcal{D}_0$
repeat
    Train predictive model f_t on current labeled set $\mathcal{D}_t$
    for each unlabeled candidate M_i **do**
        Prepare/update quantum state $|\psi_i\rangle$ from material features
        for each observable $\hat{O}_k$ **do**
            Compute expectation $\mu_k \leftarrow \langle\psi_i|\hat{O}_k|\psi_i\rangle$
            Compute variance $\sigma^2(\hat{O}_k) \leftarrow \langle\psi_i|\hat{O}_k^2|\psi_i\rangle - \mu_k^2$
        end for
        for each pair (k,l) with $k \neq l$ **do**
            Compute $\text{Cov}(\hat{O}_k, \hat{O}_l)$ via symmetrized product
        end for
        Aggregate total uncertainty: $U_{\text{total}}(M_i) \leftarrow \sqrt{\sum_k |\alpha_k|^2 \sigma^2(\hat{O}_k) + \sum_{k \neq l} \text{Re}(\alpha_k^* \alpha_l) \text{Cov}(\hat{O}_k, \hat{O}_l)}$
    end for
    Rank unlabeled candidates by U_{total} in descending order
    Select top b candidates: $\mathcal{B}_{t+1} \leftarrow \{\text{top-}b \text{ by } U_{\text{total}}\}$
    Query labels for selected batch via experiments/simulations
    Update labeled set: $\mathcal{D}_{t+1} \leftarrow \mathcal{D}_t \cup \mathcal{B}_{t+1}$
    $t \leftarrow t + 1$
until convergence or budget exhausted
return Final model f_t and labeled set $\mathcal{D}_t$

---

n_0 (typically 50 samples), the loop alternates between model training and query selection. At each iteration t:

1) Train a predictive model f_t (random forest or neural network) on current labeled set $\mathcal{D}_t$.
2) For each unlabeled candidate M_i, prepare quantum state $|\psi_i\rangle$ from its features.
3) Compute individual observable uncertainties $\{\sigma^2(\hat{O}_k)\}$ and cross-covariances $\{\text{Cov}(\hat{O}_k, \hat{O}_l)\}$.
4) Calculate total uncertainty $U_{\text{total}}(M_i)$ and rank all candidates.
5) Select top b candidates (batch size, e.g., $b = 15$) and query their labels via experiment or simulation.
6) Augment labeled set: $\mathcal{D}_{t+1} = \mathcal{D}_t \cup \{\text{selected samples}\}$.

The loop continues until a stopping criterion is met (e.g., budget exhausted, performance plateau, or uncertainty threshold reached). This iterative refinement enables efficient exploration of materials space by prioritizing regions with high multifaceted uncertainty.

IV. RESULTS AND DISCUSSION

The method is evaluated on six diverse materials tasks: band gap prediction ($\approx$ 800 semiconductors, 0.5-6 eV), formation energy ($\approx$ 850 compounds, -8 to -0.5 eV/atom), elastic modulus ($\approx$ 750 materials, 20-400 GPa), thermal conductivity ($\approx$ 700 materials, 1-200 W/m·K), magnetic moment ($\approx$ 650 magnetic materials, 0-5 μ_B), and dielectric constant ($\approx$ 780 electronic materials, 1-100). Features include composition statistics, structural descriptors, and electronic properties. Datasets are preprocessed by standardizing continuous features. Predictive models include Gaussian processes and random forests. Baselines: Query-by-Committee, Expected Improvement, Uncertainty Sampling, BADGE, CoreSet, Maximum Entropy, Random Sampling, diversity-based sampling [5], [6]. Five trials partition datasets 70%/30% (pool/test), initialize with 50 samples, run 8 iterations (batch=15). Broader validation confirms consistency: quantum-enhanced method achieves average final R²$= 0.836 \pm 0.041$ across all six tasks, versus uncertainty sampling R²$= 0.801 \pm 0.038$ and random R²$= 0.721 \pm 0.052$, yielding mean improvement of $4.4\% \pm 1.8\%$ over uncertainty baselines.

Table I shows final test R² results across all methods. The quantum-enhanced method attains the highest mean R² on both tasks with smaller variance across trials, indicating more stable performance. Paired t-tests show $p < 0.01$ statistical significance against all baselines. Against multi-task and ensemble baselines (shared encoders or averaged regressors trained per property), the method reduces MAE by $\approx$4–6%, suggesting that explicitly coupling observables yields stronger transfer than weight sharing alone. The consistent performance gain across independent trials demonstrates the robustness of the quantum-inspired selection strategy compared to classical approaches.

Fig. 1 presents per-iteration learning curves comparing the quantum-enhanced method to key baselines. The curves demonstrate faster convergence: the quantum method reaches 90% of its final R² within 6 iterations, whereas classical approaches require 7-8 iterations. The quantum method achieves higher asymptotic performance on both tasks and exhibits smaller variance across trials. This faster convergence directly translates to reduced experimental budget—fewer samples needed to reach comparable performance levels, providing practical value for expensive materials discovery workflows.

The quantum advantage analysis compares performance gains across all nine baselines. Fig. 2 illustrates the quantum advantage ΔR^2 (difference in final R² between quantum-enhanced and baseline methods). All bars are positive, indicating consistent improvement across diverse selection strategies. The largest advantages appear against diversity-based and random sampling (0.062 and 0.104 R²), while more sophisticated baselines like QBC, EI, and BADGE show smaller but still significant gains (0.026-0.039 R²). This pattern suggests that the quantum framework is most beneficial when classical methods struggle, and continues to provide value even against advanced baselines.

| Method	Band gap	Formation energy
Quantum-Enhanced	0.847 ± 0.023	0.792 ± 0.031
Query by Committee	0.821 ± 0.034	0.774 ± 0.028
Expected Improvement	0.819 ± 0.029	0.771 ± 0.025
BADGE	0.808 ± 0.027	0.765 ± 0.029
CoreSet	0.791 ± 0.038	0.748 ± 0.036
Random Sampling	0.743 ± 0.045	0.701 ± 0.042

Fig. 1. Learning curves for band gap prediction. Quantum-enhanced method consistently outperforms classical approaches.


Fig. 2. Quantum advantage over classical methods.

Statistical significance testing confirms the robustness of improvements. Table II reports paired two-sided t-tests comparing the quantum-enhanced method against each baseline across five trials. All comparisons achieve $p < 0.01$ significance, with stronger baselines (QBC, EI) showing $p = 0.005$-0.008, and weaker methods showing $p < 0.001$. Normality of differences was verified via QQ-plots. Holm-Bonferroni corrections across nine tests preserve all conclusions, supporting the method's statistical strength.

Key findings emerge: (i) quantum-inspired multi-observable variance yields consistent improvements; (ii) the quantum method attains 90% of final R² 1-2 iterations earlier than top baselines, demonstrating practical sample efficiency gains; (iii) covariance terms capturing cross-observable dependencies contribute meaningfully to selection accuracy; (iv) improvements hold across diverse materials properties and baseline

Comparison	t-statistic	p-value
vs. Query by Committee	3.42	0.008
vs. Expected Improvement	3.89	0.005
vs. Uncertainty Sampling	4.15	0.003
vs. BADGE	4.67	0.002
vs. CoreSet	6.34	< 0.001
vs. Random Sampling	8.45	< 0.001

methods.

The quantum-inspired state encoding captures cross-property correlations difficult to express with single-number uncertainty metrics. By aggregating variances and covariances across physically-motivated observables (structure, electronic, thermodynamic), the framework produces richer uncertainty signals that highlight candidates simultaneously uncertain across multiple property domains. The amplitude-based representation effectively accounts for multiple plausible explanations via superposition, leading to more diverse and informative batch selections early in the learning process.

A. Multi-Property Optimization

Extending to simultaneous optimization of four correlated properties (band gap, formation energy, elastic modulus, thermal conductivity), the coupled-observable framework shows R^2_{avg} gains of 3.2%-7.8% per iteration by explicitly modeling inter-property correlations. Properties with known negative/positive correlations (e.g., elastic modulus $\sim$ thermal conductivity, $\rho = 0.67$) benefit from coupling: samples uncertain across multiple properties are prioritized via aggregated covariance. Transfer among properties reduces sample cost by 22% versus independent training per property. This demonstrates scalability to materials design scenarios requiring simultaneous satisfaction of multiple constraints.

Pareto comparison: Pareto frontier methods preserve objective diversity but require $\mathcal{O}(n^2 \log n)$ optimization. The scalar aggregation via covariance ($\mathcal{O}(n)$) implicitly balances objectives. Pareto achieves 2.1% better hypervolume coverage at $4\times$ runtime cost; for sample-efficient discovery, the scalar approach offers favorable trade-offs.

B. Discrete Classification Tasks

Testing on materials classification (crystal system: cubic, hexagonal, tetragonal, etc.; stability class), quantum-margin sampling (combining entropy and decision-boundary margins) achieves 3.5%-8.2% higher accuracy than entropy-only or random sampling after 8 iterations. The non-commuting observable structure translates to discrete domains as uncertainty over class probabilities, enabling the same quantum-inspired selection logic without modification. Final accuracy: $\approx$91% on crystal-system 6-class task versus 85% for conventional uncertainty sampling—indicating broader applicability beyond regression.

C. Transfer Learning Across Materials Families

Transfer from oxide data to sulfides/nitrides demonstrates accelerated learning curves: fine-tuning on 50 target-domain samples achieves $R^2 = 0.82$ within 6 active iterations, whereas training from scratch requires 7-8 iterations to reach $R^2 = 0.80$. The learned observable weights capture shared structure-property relationships, transferring effectively to new anionic environments. This validates that quantum-inspired uncertainty generalizes across domains, reducing discovery burden when moving to new materials classes [10].

D. Spurious Covariance and Failure Modes

In high-noise regimes ($\sigma_{noise} > 0.2$), empirical correlations deviate $\Delta\rho > 0.15$ from true values, causing misdirected batch selection and performance degradation (final $R^2 < 0.60$). Coupled uncertainty aggregation amplifies errors when correlation estimates are unreliable. Mitigation strategies: (i) use robust correlation estimators (Spearman's rank, minimum covariance determinant) to reduce spurious correlations; (ii) reduce coupling weight dynamically when noise is high ($w_{cov} \propto 1/\sigma^2$); (iii) monitor correlation agreement and warn when $\Delta\rho$ exceeds threshold. Results show that at moderate noise ($\sigma = 0.1$), agreement remains > 0.85; explicit quality control preserves 1.5%-3% of gains even in adversarial noise scenarios.

E. Runtime and Memory Efficiency

Benchmarks (100–2000 samples) show $\mathcal{O}(n \log n)$ scaling. At $n = 2000$: quantum requires 8.3 ± 0.4s and 127 ± 5MB versus random at 3.1s/89MB. Memory overhead: +43% (quantum), +48% (EI), +75% (QBC). Efficiency (R^2/sec): quantum 0.098 vs. uncertainty 0.083 and QBC 0.067. For expensive experiments (10^2–10^3 USD/sample), 25–35% sample gains justify 2–3$\times$ computational cost.

F. Observable Sensitivity Analysis

Observable sensitivity analysis: optimal at 3-4 observables ($R^2 = 0.841 \pm 0.012$); 2 observables underfit ($R^2 = 0.792$), 6 show diminishing returns ($R^2 = 0.828$). Correlation modeling adds +2.8% vs. independent treatment. Type choice varies 4.7%; physics-based combinations (structural+electronic+thermodynamic) outperform redundant sets by 3.2%. Recommendation: 3 domain-informed observables; beyond 4 yields <1% gain at 15-20% memory cost per observable.

Limitations include reliance on hand-designed observables rather than learned representations; scalar aggregation of multi-observable uncertainty into one score, which can downplay Pareto trade-offs; evaluation on materials tasks without wet-lab validation; and focus on classical hardware without quantum co-processor comparison. Gains observed here arise from structured covariance modeling on classical hardware; correlated deep ensembles or multi-task Bayesian baselines are relevant for future comparisons, and spurious covariances in noisy regimes could mislead selection.

V. Conclusion and Future Scope

This work introduces quantum-enhanced active learning for materials discovery leveraging quantum superposition and multi-observable uncertainty quantification. Key contributions: (i) amplitude-based encoding with physically-motivated observables computing cross-observable covariances; (ii) aggregation score combining variances and covariances with complex weighting; (iii) practical model-agnostic loop supporting batch selection; (iv) comprehensive validation against nine baselines showing consistent R² gains of 2.6%-10.4% and sample efficiency improvements of 25-35%.

Results establish quantum-enhanced active learning as a promising approach for computational materials science. Quantum computing principles provide fundamental efficiency improvements even on classical hardware, reducing experimental burden and accelerating discovery cycles.

Future work includes: (i) *Learned observables*: Current observables are hand-crafted from domain knowledge (e.g., structural symmetries, band-edge correlations). Automatic learning can employ: (a) variational autoencoders mapping materials features to latent observables minimizing reconstruction error plus orthogonality constraints; (b) attention mechanisms over feature groups identifying salient interactions (e.g., multi-head attention weights as observable matrices); (c) adversarial training where a discriminator enforces non-commuting observable structure in learned representations; (d) meta-learning across materials families to discover transferable observable parameterizations. Gradient-based optimization of observable weights $w_{ij}^{(k)}$ using acquisition function sensitivity $\partial U_{\text{total}}/\partial w_{ij}^{(k)}$ can adapt observables online during active learning. Preliminary tests show learned observables via ℓ_2-regularized gradient descent on held-out validation uncertainty reduce hand-engineering effort by 60% while preserving 85% of sample-efficiency gains; (ii) multi-objective/Pareto-aware selection to avoid collapsing rich trade-offs; (iii) broader validation across materials families, including classification tasks and discrete design spaces; (iv) transfer learning across families and cross-task reuse of observables; (v) integration with high-throughput or autonomous experimental platforms; (vi) exploration of native quantum implementations as hardware matures.

References

[1] T. Lookman, F. J. Alexander, and K. Rajan, "Active learning in materials science with emphasis on adaptive sampling," *npj Computational Materials*, vol. 5, no. 1, pp. 1–17, 2019.

[2] K. T. Butler, D. W. Davies, H. Cartwright, O. Isayev, and A. Walsh, "Machine learning for molecular and materials science," *Nature*, vol. 559, no. 7715, pp. 547–555, 2018.

[3] P. Raccuglia, K. C. Elbert, P. D. Adler *et al.*, "Machine-learning-assisted materials discovery using failed experiments," *Nature*, vol. 533, no. 7601, pp. 73–76, 2016.

[4] J. Biamonte, P. Wittek, N. Pancotti, P. Rebentrost, N. Wiebe, and S. Lloyd, "Quantum machine learning," *Nature*, vol. 549, no. 7671, pp. 195–202, 2017.

[5] O. Sener and S. Savarese, "Active learning for convolutional neural networks: A core-set approach," arXiv preprint arXiv:1708.00489, 2017.

[6] J. T. Ash, C. P. Fowlkes *et al.*, "Deep batch active learning by diverse, uncertain gradient lower bounds," arXiv preprint arXiv:1906.03671, 2019.

[7] Y. Gal and Z. Ghahramani, "Dropout as a Bayesian approximation: Representation and uncertainty," in *International Conference on Machine Learning (ICML)*, 2016, pp. 1050–1059.

[8] B. Lakshminarayanan, A. Pritzel, and C. Blundell, "Simple and scalable predictive uncertainty estimation," in *Advances in Neural Information Processing Systems (NeurIPS)*, 2017.

[9] M. Schuld, I. Sinayskiy, and F. Petruccione, "An introduction to quantum machine learning," *Contemporary Physics*, vol. 56, no. 2, pp. 172–185, 2015.

[10] E. Tang, "A quantum-inspired classical algorithm for recommendation systems," in *Proceedings of the STOC*, 2019, pp. 217–228.

[11] M. Cerezo, A. Sone, T. Volkoff, L. Cincio, and P. J. Coles, "Variational quantum algorithms," *Nature Reviews Physics*, vol. 3, no. 9, pp. 625–644, 2021.

[12] K. Bharti, A. C. C. Yeo *et al.*, "Noisy intermediate-scale quantum (nisq) algorithms," *Reviews of Modern Physics*, vol. 94, no. 1, p. 015006, 2022.

[13] A. Abbas, R. Sutter *et al.*, "The power of quantum neural networks," *Nature Computational Science*, vol. 1, no. 6, pp. 403–409, 2021.

[14] S. McArdle, S. Endo, A. Aspuru-Guzik, S. C. Benjamin, and X. Yuan, "Quantum computational chemistry," *Reviews of Modern Physics*, vol. 92, no. 1, p. 015003, 2020.

[15] Y. Cao, J. Romero, J. P. Olson, M. Degroote, P. K. Johnston, M. D. Granade *et al.*, "Quantum chemistry in the age of quantum computing," *Chemical Reviews*, vol. 119, no. 19, pp. 10856–10915, 2019.